

Kubernetes Certified Application Developer (CKAD)

Guide v1.0beta

Author: Sujay Bhowmick

Email: sujaybhowmick at gmail dot com

References:

- from and others [dgkanatsios](#)

Kubernetes Objects

Deployment Object

Creating a Deployment

```
# deployment-definition.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  namespace: default
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.7.9
          ports:
            - containerPort: 80
```

```
$ kubectl create -f deployment-definition.yaml
```

List deployments

```
$ kubectl get deployments # default namespace
$ kubectl get deployments --namespace=<namespace-name>
$ kubectl get deployment <deployment-name> # default namespace
$ kubectl get deployment <deployment-name> --namespace=<namespace-name>
```

Deployment details

```
$ kubectl describe deployment <deployment-name> # default namespace
$ kubectl get deployment <deployment-name> --namespace=<namespace-name>
```

Delete deployment

```
$ kubectl delete deployment <deployment-name>
$ kubectl delete deployment <deployment-name> --namespace=<namespace-name>
```

Update a deployment

```
$ kubectl apply -f deployment-definition.yaml
```

Pod Object

Creating a Pod

In general, users shouldn't need to create pods directly. They should almost always use controllers even for singletons, for example, [Deployments](#). Controllers provide self-healing with a cluster scope, as well as replication and rollout management. Controllers use a Pod Template that you provide to create the Pods for which it is responsible.

```
# pod-definition.yaml - POD Template
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  namespace: default
  labels:
    app: myapp
spec:
  containers:
  - name: myapp-container
    image: busybox
    command: ['sh', '-c', 'echo Hello Kubernetes! && sleep 3600']
```

```
$ kubectl create -f pod-definition.yaml
```

List Pods

```
$ kubectl get pods # default namespace
$ kubectl get pods --namespace=<namespace-name>
```

Pod details

```
$ kubectl describe pod <pod-name> # default namespace
$ kubectl describe pod <pod-name> --namespace=<namespace-name>
```

Delete Pod

```
$ kubectl delete pod <pod-name> # default namespace
$ kubectl delete pod <pod-name> --namespace=<namespace-name>
```

ReplicaSets Object

Create a ReplicaSet

It is recommend using Deployments instead of directly using ReplicaSets, unless you require custom update orchestration or don't require updates at all.

we recommend using Deployments instead of directly using ReplicaSets, unless you require custom update orchestration or don't require updates at all.

```
# replicaset-definition.yaml - ReplicaSet Template
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: frontend
  namespace: default
  labels:
    app: guestbook
    tier: frontend
spec:
  # modify replicas according to your case
  replicas: 3
  selector:
    matchLabels:
      tier: frontend
  template:
    metadata:
      labels:
        tier: frontend
    spec:
      containers:
        - name: php-redis
          image: gcr.io/google_samples/gb-frontend:v3
```

```
$ kubectl create -f replicaset-definition.yaml
```

List ReplicaSets

```
$ kubectl get replicaset # default namespace
$ kubectl get replicaset --namespace=<namespace-name>
```

ReplicaSet details

```
$ kubectl describe replicaset <replicaset-name> # default namespace
$ kubectl describe replicaset <replicaset-name> --namespace=<namespace-name>
```

Delete ReplicaSet

```
$ kubectl delete replicaset <replicaset-name> # default namespace
$ kubectl delete replicaset <replicaset-name> --namespace=<namespace-name>
```

Namespace Object

Create a namespace

```
# namespace-definition.yaml - Namespace template
apiVersion: v1
kind: Namespace
metadata:
  name: development-namespace
```

```
$ kubectl create -f namespace-definition.yaml
```

List Namespace

```
$ kubectl get namespaces
```

Namespace details

```
$ kubectl describe namespace <namespace-name>
```

Delete Namespace

Warning: This deletes *everything* under the namespace!

```
$ kubectl delete namespace <namespace-name>
```

ConfigMap Object

Create a ConfigMap

```
# creating configmap from properties files in a directory (key=value)
$ kubectl create configmap <configmap-name> --from-file=configMap/
# creating configmap from env files in a directory (VAR=VALUE)
$ kubectl create configmap <configmap-name> --from-env-file=configMap/
# creating configmap from literal value
$ kubectl create configmap <configmap-name> --from-literal=VAR=VALUE
```

List ConfigMaps

```
$ kubectl get configmaps # default namespace
$ kubectl get configmaps --namespace=<namespace-name>
```

ConfigMap details

```
$ kubectl describe configmap <config-map-name> # default namespace
$ kubectl describe configmap <config-map-name> --namespace=<namespace-name>
```

Delete ConfigMap

```
$ kubectl delete configmap <config-map-name> # default namespace
$ kubectl delete configmap <config-map-name> --namespace=<namespace-name>
```

Using ConfigMaps

```
# example usage of configmaps in Deployments
```

Secrets Object

Create a Secret

```
# secrets-definition.yaml
apiVersion: v1
kind: Secret
metadata:
  name: mysecret
type: Opaque
data:
  username: YWRtaW4=
  password: MWYyZDFlMmU2N2Rm
```

```
$ kubectl create -f secrets-definition.yaml
# creating secret from files
$ kubectl create secret generic db-user-pass --from-file=./username.txt --from-file=./password.txt
# creating secret using literal
$ kubectl create secret generic prod-db-secret --from-literal=username=produser --from-literal=password=Y4nys7f1l
```

List Secrets

```
$ kubectl get secrets # default namespace
$ kubectl get secrets --namespace=<namespace-name>
```

Secret details

```
$ kubectl describe secret <secret-name> # default namespace
$ kubectl describe secret <secret-name> --namespace=<namespace-name>
```

Delete a Secret

```
$ kubectl delete secret <secret-name> # default namespace
$ kubectl delete secret <secret-name> --namespace=namespace
```

ServiceAccount Object

Create a ServiceAccount

```
# service-account-definition.yaml
apiVersion: v1
kind: ServiceAccount
metadata:
  name: secret-access-sa
```

```
$ kubectl create -f service-account-definition.yaml
```

List ServiceAccounts

```
$ kubectl get serviceaccounts # default namespace
$ kubectl get serviceaccounts --namespace=<namespace-name>
```

ServiceAccount details

```
$ kubectl describe serviceaccount <service-account-name> # default namespace
$ kubectl describe serviceaccount <service-account-name> --namespace=<namespace-name>
```

Delete a ServiceAccount

```
$ kubectl delete serviceaccount <service-account-name> # default namespace
$ kubectl delete serviceaccount <service-account-name> --namespace=<namespace-name>
```

Service Object

Create a Service

```
# service-definition.yaml
apiVersion: v1
kind: Service
metadata:
  name: my-service
  namespace: default
spec:
  selector:
    app: MyApp
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```

```
$ kubectl create -f service-definition.yaml
```

List Services

```
$ kubectl get services # default namespace
$ kubectl get services --namespace=<namespace-name>
```

Service details

```
$ kubectl describe service <service-name> # default namespace
$ kubectl describe service <service-name> --namespace=<namespace-name>
```

Delete a Service


```
$ kubectl delete service <service-name> # default namespace
$ kubectl delete service <service-name> --namespace=<namespace-name>
```

ResourceQuota Object

Create a Quota

```
# resource-quota-definition.yaml
apiVersion: v1
kind: ResourceQuota
metadata:
  name: myresourcequota
  namespace: default
spec:
  hard:
    cpu: "1"
    memory: 1G
    pods: "2"
```

```
$ kubectl create -f resource-quota-definition.yaml
# or
$ kubectl create quota myrq --hard=cpu=1,memory=1G,pods=2 --dry-run -o yaml
```

List ResourceQuota

```
$ kubectl get quota <resource-quota-name> # default namespace
$ kubectl get quota <resource-quota-name> --namespace=<namespace-name>
```

ResourceQuota details

```
$ kubectl describe quota <resource-quota-name> # default namespace
$ kubectl describe quota <resource-quota-name> --namespace=<namespace-name>
```

Delete a ResourceQuota

```
$ kubectl delete quota <resource-quota-name> # default namespace
$ kubectl delete quota <resource-quota-name> --namespace=<namespace-name>
```

Configuring Volumes in Kubernetes

Configure a Volume

```
# pod-volume-aws-definition.yaml for AWS block store
apiVersion: v1
kind: Pod
metadata:
  name: rand-num-gen
  namespace: default
spec:
  containers:
    - name: alpine
      image: alpine
      command: ["/bin/sh" "-c"]
      args: ["shuf -i 0-100 -n 1 >> /opt/number.out;"]
      volumeMounts:
        - mountPath: /opt
          name: data-volume

  volumes:
    - name: data-volume
      awsElasticBlockStore:
        volumeID: <volume-id>
        fsType: ext4
```

Do not use the node directory in a cluster setup, only for single node.

```
# pod-volume-definition.yaml for node directory
apiVersion: v1
kind: Pod
metadata:
  name: rand-num-gen
  namespace: default
spec:
  containers:
    - name: alpine
      image: alpine
      command: ["/bin/sh", "-c"]
      args: ["shuf -i 0-100 -n 1 >> /opt/number.out;"]
      volumeMounts:
        - mountPath: /opt
          name: data-volume

  volumes:
    - name: data-volume
      hostPath:
        path: /data
        type: Directory
```

Exercise

Core Concepts (13%)

1. Create a namespace called 'mynamespace' and a pod with image *nginx* called *nginx* on this namespace

```
# Answer - Namespace
apiVersion: v1
kind: Namespace
metadata:
  name: mynamespace
```

```
# Answer - Pod
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: mynamespace
spec:
  containers:
  - name: nginx
    image: nginx
```

2. Create a busybox pod (using kubectl command) that runs the command "env". Run it and see the output

```
# $ kubectl run busybox --image=busybox --restart=Never --dry-run -o yaml --
command -- env > envpod.yaml
# $ kubectl logs busybox
apiVersion: v1
kind: Pod
metadata:
  name: busybox
  namespace: default
spec:
  containers:
  - name: busybox
    image: busybox
    command:
    - "env"
```

3. Get the YAML for a new namespace called 'myns' without creating it

```
# Get the YAML for a new namespace called 'myns' without creating it
# $ kubectl create namespace myns --dry-run -o yaml
apiVersion: v1
kind: Namespace
metadata:
  creationTimestamp: null
  name: myns
spec: {}
status: {}
```

4. Get pods on all namespaces

```
$ kubectl get pods --all-namespaces
```

5. Create a pod with image nginx called nginx and allow traffic on port 80

```
# Create a pod with image nginx called nginx and allow traffic on port 80
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: default
spec:
  containers:
  - name: nginx
    image: nginx
    ports:
    - containerPort: 80
```

6. Get the pod's ip, use a temp busybox image to wget its '/'

```
$ kubectl get po -o wide # get the IP, will be something like '10.1.1.131'
# create a temp busybox pod
$ kubectl run busybox --image=busybox --rm -it --restart=Never -- sh
# run wget on specified IP:Port
wget -O- 10.1.1.131:80
exit

# Or alternatively
# Get IP of the nginx pod
NGINX_IP=$(kubectl get pod nginx -o jsonpath='{.status.podIP}')
# create a temp busybox pod
```

```
$ kubectl run busybox --image=busybox --env="NGINX_IP=$NGINX_IP" --rm -it --
restart=Never -- sh
# run wget on specified IP:Port
wget -O- $NGINX_IP:80
exit
```

7. Get this pod's YAML without cluster specific information

```
$ kubectl get pod nginx -o yaml --export
```

8. Get information about the pod, including details about potential issues (e.g. pod hasn't started)

```
$ kubectl describe pod <pod-name>
```

9. Get pod logs

```
$ kubectl logs <pod-name>
```

10. If pod crashed and restarted, get logs about the previous instance

```
$ kubectl logs <pod-name> -p
```

11. Connect to the nginx pod

```
$ kubectl exec -it <pod-name> -- /bin/sh
```

12. Create a busybox pod that echoes 'hello world' and then exits

```
# Create a busybox pod that echoes 'hello world' and then exits
apiVersion: v1
kind: Pod
metadata:
  name: busybox
  namespace: default
spec:
  containers:
  - name: busybox
    image: busybox
    command:
      - "echo"
    args:
      - "hello world"
```

14. Do the same, but have the pod deleted automatically when it's completed

```
$ kubectl run busybox --image=busybox -it --rm --restart=Never -- /bin/sh -c  
'echo hello world'  
$ kubectl get po
```

15. Create an nginx pod and set an env value as 'var1=val1'. Check the env value existence within the pod

```
$ kubectl run nginx --image=nginx --env=var1=val1  
$ kubectl exec -it nginx -- env
```

16. Get the YAML for a new ResourceQuota called 'myrq' without creating it

```
$ kubectl create quota myrq --hard=cpu=1,memory=1G,pods=2 --dry-run -o yaml
```

Pod Design (20%)

Labels and Annotations

1. Create 3 pods with names nginx1,nginx2,nginx3. All of them should have the label app=v1

```
$ kubectl run nginx1 --image=nginx --labels=app=v1  
$ kubectl run nginx2 --image=nginx --labels=app=v1  
$ kubectl run nginx3 --image=nginx --labels=app=v1
```

2. Show all labels of the pods

```
$ kubectl get pods --show-labels
```

3. Change the labels of pod 'nginx2' to be app=v2

```
$ kubectl label pod nginx2 app=v2 --overwrite
```

4. Get the label 'app' for the pods

```
$ kubectl get pod -L app
```

5. Get only the 'app=v2' pods

```
$ kubectl get pod -l app=v2
# or
$ kubectl get pod -l 'app in (v2)'
```

6. Remove the 'app' label from the pods we created before

```
$ kubectl label pod -lapp app-
# or
$ kubectl label pod nginx{1..3} app-
```

7. Create a pod that will be deployed to a Node that has the label 'accelerator=nvidia-tesla-p100'

```
# Create a pod that will be deployed to a Node that has the label
'accelerator=nvidia-tesla-p100'
apiVersion: v1
kind: Pod
metadata:
  name: fastpod
  namespace: default
spec:
  containers:
    - name: fastpod
      image: nginx
  nodeSelector:
    accelerator: nvidia-tesla-p100
```

8. Annotate pods nginx1, nginx2, nginx3 with "description='my description'" value

```
$ kubectl annotate nginx1, nginx2, nginx3 description='my description'
```

9. Check the annotations for pod nginx1

```
$ kubectl describe pod nginx1 | grep -i annotations
```

10. Remove the annotations for these three pods

```
$ kubectl annotate pod nginx{1..3} description-
```

Deployments

1. Create a deployment with image nginx:1.7.8, called nginx, having 2 replicas, defining port 80 as the port that this container exposes (don't create a service for this deployment)

```

# Create a deployment with image nginx:1.7.8, called nginx, having 2
replicas,
# defining port 80 as the port that this container exposes
# (don't create a service for this deployment)
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx
  namespace: default
spec:
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      name: nginx
      namespace: default
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.7.8
          ports:
            - containerPort: 80
  replicas: 2

```

2. View the YAML of this deployment

```
$ kubectl get deployment nginx --export -o yaml
```

3. View the YAML of the ReplicaSet created

```

$ kubectl get replicaset -l app=nginx
# or
$ kubectl get rs -l app=nginx

```

4. Get the YAML for one of the pods

```

$ kubectl get po # list of all pods
$ kubectl get po nginx-5746858fb4-l6qfb -o yaml --export

```

5. Check how the deployment rollout is going


```
$ kubectl rollout status deployment nginx
```

6. Update the nginx image to nginx:1.7.9

```
$ kubectl edit deployment nginx  
# or  
$ kubectl set image deploy nginx nginx=nginx:1.7.9
```

7. Undo the latest rollout and verify that new pods have the old image (nginx:1.7.8)

```
$ kubectl rollout undo deploy nginx
```

8. Do an on purpose update of the deployment with a wrong image nginx:1.91

```
$ kubectl set image deploy nginx nginx=nginx:1.19
```

9. Verify that something's wrong with the rollout

```
$ kubectl rollout status deploy nginx
```

10. Return the deployment to the second revision (number 2) and verify the image is nginx:1.7.9

```
$ kubectl rollout undo deploy nginx --to-revision=4
```

11. Check the details of the third revision rollout (number 4)

```
$ kubectl rollout history deploy nginx --revision=4
```

12. Scale the deployment to 5 replicas

```
$ kubectl scale --current-replicas=2 --replicas=5 deploy nginx
```

13. Autoscale the deployment, pods between 5 and 10, targetting CPU utilization at 80%

```
$ kubectl autoscale deploy nginx --min=5 --max=10 --cpu-percent=80
```

14. Pause the rollout of the deployment

```
$ kubectl rollout pause deploy nginx
```

15. Update the image to nginx:1.9.1 and check that there's nothing going on, since we paused the rollout

```
$ kubectl set image deploy nginx nginx=nginx:1.9.1
$ kubectl rollout history deploy nginx
```

16. Resume the rollout and check that the nginx:1.9.1 image has been applied

```
$ kubectl rollout resume deploy nginx
$ kubectl rollout history deploy nginx
```

17. Delete the deployment and the horizontal pod autoscaler you created

```
$ kubectl delete deploy nginx
$ kubectl delete hpa nginx
```

Jobs

1. Create a job with image perl that runs default command with arguments "perl -Mbignum=bpi -wle 'print bpi(2000)'"

```
apiVersion: batch/v1
kind: Job
metadata:
  name: perl
  namespace: default
spec:
  template:
    spec:
      containers:
        - name: perl
          image: perl
          command: ["perl", "-Mbignum=bpi", "-wle", "print bpi(2000)"]
          restartPolicy: Never
```

2. Wait till it's done, get the output

```
$ kubectl logs perl-h9sbz
```

3. Create a job with the image busybox that executes the command 'echo hello;sleep 30;echo world'

```
apiVersion: batch/v1
kind: Job
```

```

metadata:
  name: busybox-job
  namespace: default
spec:
  template:
    spec:
      containers:
      - name: busybox-job
        image: busybox
        command: ["/bin/sh"]
        args: ["-c", "echo hello; sleep 30; echo world"]
        restartPolicy: OnFailure

```

4. See the status of the job, describe it and see the logs

```

$ kubectl get job busybox-job
$ kubectl describe job busybox-job
$ kubectl logs job/busybox-job

```

5. Create the same job, make it run 5 times, one after the other. Verify its status and delete it

```

apiVersion: batch/v1
kind: Job
metadata:
  name: busybox-job
  namespace: default
spec:
  completions: 5
  template:
    spec:
      containers:
      - name: busybox-job
        image: busybox
        command: ["/bin/sh"]
        args: ["-c", "echo hello; sleep 30; echo world"]
        restartPolicy: OnFailure

```

6. Create the same job, but make it run 5 parallel times

```

apiVersion: batch/v1
kind: Job
metadata:
  name: busybox-job
  namespace: default
spec:

```

```

completions: 5
parallelism: 5
template:
  spec:
    containers:
      - name: busybox-job
        image: busybox
        command: ["/bin/sh"]
        args: ["-c", "echo hello; sleep 30; echo world"]
        restartPolicy: OnFailure

```

CronJobs

1. Create a cron job with image busybox that runs on a schedule of "`*/* * * * *`" and writes 'date; echo Hello from the Kubernetes cluster' to standard output

```

apiVersion: batch/v1beta1
kind: CronJob
metadata:
  name: busybox-cronjob
  namespace: default
spec:
  schedule: "*/1 * * * *"
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: busybox-cronjob
              image: busybox
              command: ["/bin/sh"]
              args: ["-c", "date; echo Hello from Kubernetes cluster"]
              restartPolicy: OnFailure

```

2. See its logs and delete it

```

$ kubectl get cj
$ kubectl get jobs --watch
$ kubectl get po --show-labels
$ kubectl logs busybox-cronjob-1556066760-smzn7
$ kubectl delete cj busybox-cronjob

```

Configuration (18%)

ConfigMaps

1. Create a configmap named config with values foo=lala,foo2=lolo

```
$ kubectl create configmap config --from-env-file=app.properties
# or
$ kubectl create configmap config --from-literal=foo=lala --from-
literal=foo2=lolo
```

2. Display its values

```
$ kubectl get cm config -o yaml --export
# or
$ kubectl describe cm config
```

3. Create and display a configmap from a file, giving the key 'special'

```
$ kubectl create configmap config --from-file=special=app.properties
```

4. Create a configMap called 'options' with the value var5=val5. Create a new nginx pod that loads the value from variable 'var5' in an env variable called 'option'

```
$ kubectl create configmap options --from-literal=var5=val5
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: default
spec:
  containers:
    - name: nginx
      image: nginx
      env:
        - name: option
          valueFrom:
            configMapKeyRef:
              name: options
              key: var5
      restartPolicy: Never
```

```
# to check env variables
$ kubectl exec nginx -it -- env
```

5. Create a configMap 'anotherone' with values 'var6=val6', 'var7=val7'. Load this configMap as env variables into a new nginx pod

```
$ kubectl create configmap anotherone --from-literal=var6=val6 --from-literal=var7=val7
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: default
spec:
  containers:
    - name: nginx
      image: nginx
      envFrom:
        - configMapRef:
            name: anotherone
      restartPolicy: Never
```

```
# to check env variables
$ kubectl exec nginx -it -- env
```

6. Create a configMap 'cmvolume' with values 'var8=val8', 'var9=val9'. Load this as a volume inside an nginx pod on path '/etc/lala'. Create the pod and 'ls' into the '/etc/lala' directory.

```
$ kubectl create cm cmvolume --from-literal=var8=val8 --from-literal=var9=val9
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: default
spec:
  volumes:
    - name: config-volume
      configMap:
        name: cmvolume
  containers:
    - name: nginx
      image: nginx
      volumeMounts:
```

```
- mountPath: /etc/lala
  name: config-volume
restartPolicy: Never
```

```
# to check
$ kubectl exec nginx -it -- /bin/sh
# cd to /etc/lala and you should see two file var8 and var 9
$ cat var8 # will show value val8
```

Security Context

1. Create the YAML for an nginx pod that runs with the UID 101. No need to create the pod

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: default
spec:
  securityContext:
    runAsUser: 101
  containers:
    - name: nginx
      image: nginx
      restartPolicy: Never
```

2. Create the YAML for an nginx pod that has the capabilities "NET_ADMIN", "SYS_TIME" added on its single container

```
apiVersion: v1
kind: Pod
metadata:
  namespace: default
  name: nginx
spec:
  containers:
    - name: nginx
      image: nginx
      securityContext:
        capabilities:
          add: ["NET_ADMIN", "SYS_TIME"]
      restartPolicy: Never
```

Requests and limits

1. Create an nginx pod with requests cpu=100m,memory=256Mi and limits cpu=200m,memory=512Mi

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: default
spec:
  containers:
  - name: nginx
    image: nginx
    resources:
      requests:
        cpu: "100m"
        memory: "256Mi"
      limits:
        cpu: "200m"
        memory: "512Mi"
    restartPolicy: Never
```

Secrets

1. Create a secret called mysecret with the values password=mypass

```
$ kubectl create secret generic mysecret --from-literal=password=mypass
```

2. Create a secret called mysecret2 that gets key/value from a file

```
$ kubectl create secret generic mysecret2 --from-file=username
```

3. Get the value of mysecret2

```
$ kubectl get secret mysecret2 -o yaml --export
$ echo YWRtaW4K | base64 -d
```

4. Create an nginx pod that mounts the secret mysecret2 in a volume on path /etc/foo

```
apiVersion: v1
kind: Pod
metadata:
  namespace: default
  name: nginx
spec:
```



```

volumes:
  - name: myvolume
    secret:
      secretName: mysecret2
containers:
  - name: nginx
    image: nginx
    volumeMounts:
      - mountPath: /etc/foo
        name: myvolume
restartPolicy: Never

```

5. Delete the pod you just created and mount the variable 'username' from secret mysecret2 onto a new nginx pod in env variable called 'USERNAME'

```

apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: default
spec:
  containers:
    - name: nginx
      image: nginx
      env:
        - name: USERNAME
          valueFrom:
            secretKeyRef:
              key: username
              name: mysecret2
      restartPolicy: Never

```

Service Accounts

1. See all the service accounts of the cluster in all namespaces

```

$ kubectl get serviceaccounts --all-namespaces
# or
$ kubectl get sa --all-namespaces

```

2. Create a new serviceaccount called 'myuser'

```

$ kubectl create serviceaccount myuser
# or
$ kubectl create sa myuser

```

3. Create an nginx pod that uses 'myuser' as a service account

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: default
spec:
  serviceAccountName: myuser
  containers:
    - name: nginx
      image: nginx
      restartPolicy: Never
```

Observability (18%)

Liveness and Readiness probes

1. Create an nginx pod with a liveness probe that just runs the command 'ls'. Save its YAML in pod.yaml. Run it, check its probe status, delete it.

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: default
spec:
  containers:
    - name: nginx
      image: nginx
      livenessProbe:
        exec:
          command: ["ls"]
      restartPolicy: Never
```

2. Modify the pod.yaml file so that liveness probe starts kicking in after 5 seconds whereas the period of probing would be 10 seconds. Run it, check the probe, delete it.

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: default
spec:
```

```

containers:
  - name: nginx
    image: nginx
    livenessProbe:
      initialDelaySeconds: 5
      periodSeconds: 10
      exec:
        command:
          - ls
    restartPolicy: Never

```

3. Create an nginx pod (that includes port 80) with an HTTP readinessProbe on path '/' on port 80. Again, run it, check the readinessProbe, delete it.

```

apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: default
spec:
  containers:
    - name: nginx
      image: nginx
      ports:
        - containerPort: 80
      readinessProbe:
        httpGet:
          port: 80
          path: /
          initialDelaySeconds: 10
          timeoutSeconds: 3
      restartPolicy: Never

```

Logging

1. Create a busybox pod that runs 'i=0; while true; do echo "\$i: \$(date)"; i=\$((i+1)); sleep 1; done'. Check its logs

```
apiVersion: v1
kind: Pod
metadata:
  name: busybox
  namespace: default
spec:
  containers:
    - name: busybox
      image: busybox
      command: ["/bin/sh"]
      args: ["-c", "i=0; while true; do echo \"$i: $(date)\"; i=$((i+1));
sleep 1; done"]
      restartPolicy: Never
```

2. Get CPU/memory utilization for nodes (heapster must be running)

```
$ kubectl top nodes
```

Services and Networking (13%)

1. Create a pod with image nginx called nginx and expose its port 80

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  containers:
    - name: nginx
      image: nginx
      resources:
        limits:
          memory: "128Mi"
          cpu: "500m"
      ports:
        - containerPort: 80
```

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-svc
spec:
  selector:
    app: nginx
  ports:
  - port: 80
    targetPort: 80
```

2. Confirm that ClusterIP has been created. Also check endpoints

```
$ kubectl get svc nginx-svc
$ kubectl get ep
```

3. Get pod's ClusterIP, create a temp busybox pod and 'hit' that IP with wget

```
$ kubectl get svc nginx-svc
$ kubectl run busybox --rm --image=busybox -it -- sh
```

4. Convert the ClusterIP to NodePort and find the NodePort port. Hit it using Node's IP. Delete the service and the pod

```
apiVersion: v1
kind: Service
metadata:
  name: myapp
spec:
  selector:
    app: myapp
  ports:
  - port: 80
    targetPort: 80
  type: NodePort
```

5. Create a deployment called foo using image 'dgkanatsios/simpleapp' (a simple server that returns hostname) and 3 replicas. Label it as 'app=foo'. Declare that containers in this pod will accept traffic on port 8080 (do NOT create a service yet)

```
apiVersion: apps/v1
kind: Deployment
metadata:
```

```

name: foo
spec:
  replicas: 3
  selector:
    matchLabels:
      app: foo
  template:
    metadata:
      labels:
        app: foo
    spec:
      containers:
      - name: foo
        image: dgkanatsios/simpleapp
        resources:
          limits:
            memory: "128Mi"
            cpu: "500m"
        ports:
        - containerPort: 8080

```

```
$ kubectl get po -o wide # IP Column
```

6. Create a service that exposes the deployment on port 6262. Verify its existence, check the endpoints

```

apiVersion: v1
kind: Service
metadata:
  name: foo-svc
spec:
  selector:
    app: foo
  ports:
  - port: 6262
    targetPort: 8080

```

```
$ kubectl get ep
```

7. Create a temp busybox pod and connect via wget to foo service. Verify that each time there's a different hostname returned. Delete deployment and services to cleanup the cluster

```
$ kubectl run busybox --rm --image=busybox -it sh
wget -O- http://10.106.123.25:6262
```

8. Create an nginx deployment of 2 replicas, expose it via a ClusterIP service on port 80. Create a NetworkPolicy so that only pods with labels 'access: true' can access the deployment and apply it

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx
        resources:
          limits:
            memory: "128Mi"
            cpu: "500m"
        ports:
        - containerPort: 80
```

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-svc
spec:
  selector:
    app: nginx
  ports:
  - port: 80
    targetPort: 80
```

```

apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: access-nginx
spec:
  podSelector:
    matchLabels:
      app: nginx
  ingress:
    - from:
      - podSelector:
          matchLabels:
            access: 'true'

```

```

# Cannot access the pod as label access=true is missing
$ kubectl run busybox --image=busybox --restart=Never --rm -it -- wget -O-
http://IP:80
# Can access the pod as label access=true is present
$ kubectl run busybox --image=busybox --restart=Never --labels=access=true --
rm -it -- wget -O- http://IP:80

```

State Persistence (8%)

1. Create busybox pod with two containers, each one will have the image busybox and will run the 'sleep 3600' command. Make both containers mount an emptyDir at '/etc/foo'. Connect to the second busybox, write the first column of '/etc/passwd' file to '/etc/foo/passwd'. Connect to the first busybox and write '/etc/foo/passwd' file to standard output. Delete pod.

```

apiVersion: v1
kind: Pod
metadata:
  name: busybox
  labels:
    name: busybox
spec:
  volumes:
    - name: myvolume
      emptyDir: {}
  containers:
    - name: busybox
      image: busybox
      resources:
        limits:
          memory: "128Mi"

```



```

      cpu: "500m"
    ports:
      - containerPort: 80
    command: ["/bin/sh"]
    args: ["-c", "sleep 3600;"]
    volumeMounts:
      - name: myvolume
        mountPath: /etc/foo
  - name: busybox2
    image: busybox
    resources:
      limits:
        memory: "128Mi"
        cpu: "500m"
    ports:
      - containerPort: 80
    command: ["/bin/sh"]
    args: ["-c", "sleep 3600;"]
    volumeMounts:
      - name: myvolume
        mountPath: /etc/foo

```

```

# connect to container busybox2
$ kubectl exec busybox -c busybox2 -it -- /bin/sh
cat /etc/passwd | cut -f 1 -d ':' > /etc/foo/passwd
# connect to container busybox
$ kubectl exec busybox -c busybox -it -- /bin/sh
cat /etc/foo/passwd

```

2. Create a PersistentVolume of 10Gi, called 'myvolume'. Make it have accessMode of 'ReadWriteOnce' and 'ReadWriteMany', storageClassName 'normal', mounted on hostPath '/etc/foo'. Save it on pv.yaml, add it to the cluster. Show the PersistentVolumes that exist on the cluster.

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: myvolume
spec:
  storageClassName: normal
  capacity:
    storage: 10Gi
  accessModes:
    - ReadWriteOnce
    - ReadWriteMany
  hostPath:
    path: /etc/foo
```

```
$ kubectl get pv
```

3. Create a PersistentVolumeClaim for this storage class, called mypvc, a request of 4Gi and an accessMode of ReadWriteOnce and save it on pvc.yaml. Create it on the cluster. Show the PersistentVolumeClaims of the cluster. Show the PersistentVolumes of the cluster.

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mypvc
spec:
  resources:
    requests:
      storage: 4Gi
  storageClassName: normal
  accessModes:
    - ReadWriteOnce
```

```
$ kubectl get pvc
```

4. Create a busybox pod with command 'sleep 3600', save it on pod.yaml. Mount the PersistentVolumeClaim to '/etc/foo'. Connect to the 'busybox' pod, and copy the '/etc/passwd' file to '/etc/foo'.

```
apiVersion: v1
kind: Pod
metadata:
  name: busybox2
  labels:
```

```

    name: busybox2
spec:
  volumes:
    - name: myvolume
      persistentVolumeClaim:
        claimName: mypvc
  containers:
    - name: busybox2
      image: busybox
      args:
        - /bin/sh
        - -c
        - sleep 3600
      resources:
        limits:
          memory: "128Mi"
          cpu: "500m"
      ports:
        - containerPort: 80
      volumeMounts:
        - name: myvolume
          mountPath: /etc/foo

```

```

$ kubectl exec busybox -it -- /bin/sh
cp /etc/passwd /etc/foo/passwd
$ kubectl exec busybox2 -it -- /bin/sh
ls /etc/foo

```

5. Create a busybox pod with 'sleep 3600' as arguments. Copy '/etc/passwd' from the pod to your local folder.

```

$ kubectl run busybox --image=busybox -- sleep 3600
$ kubectl cp busybox:/etc/passwd .
cat passwd

```

Multi-container Pods (10%)

1. Create a Pod with two containers, both with image busybox and command "echo hello; sleep 3600". Connect to the second container and run 'ls'.

```

apiVersion: v1
kind: Pod
metadata:
  name: busybox

```

```
labels:
  name: busybox
spec:
  volumes:
    - name: myvolume
      emptyDir: {}
  containers:
    - name: busybox
      image: busybox
      resources:
        limits:
          memory: "128Mi"
          cpu: "500m"
      ports:
        - containerPort: 80
      command: ["/bin/sh"]
      args: ["-c", "echo hello; sleep 3600"]
    - name: busybox2
      image: busybox
      resources:
        limits:
          memory: "128Mi"
          cpu: "500m"
      ports:
        - containerPort: 80
      command: ["/bin/sh"]
      args: ["-c", "echo hello; sleep 3600"]
```

```
$ kubectl exec busybox -c busybox2 -it -- /bin/sh
ls
exit
```